

**VIETNAM NATIONAL UNIVERSITY-HO CHI MINH CITY**

**UNIVERSITY OF SCIENCE**

**FACULTY OF INFORMATION TECHNOLOGY**

---



# **FINAL PROJECT REPORT**

**Student : 20C14001 – Le Duong Tuan Anh**  
**Teacher : Dr. Le Thi Nhan**  
**Class : Master of Information System**  
**Subject : Data Mining**

**Ho Chi Minh city, July 2021**

# Contents

<b>1</b>	<b>General Information .....</b>	<b>3</b>
	Student Information .....	3
	Report Information .....	3
<b>2</b>	<b>Problem Statement .....</b>	<b>4</b>
<b>3</b>	<b>Data Description .....</b>	<b>5</b>
<b>4</b>	<b>Data Understanding .....</b>	<b>7</b>
4.1	Description .....	7
4.2	Overview by ScatterPlot .....	8
4.3	Statistical Information .....	9
<b>5</b>	<b>Data Preprocessing .....</b>	<b>14</b>
5.1.	Filling Zero-Values .....	14
5.1.1.	Fill by Mean method .....	15
5.1.2.	Fill by Regression Method .....	16
5.2.	Outliers Removal .....	19
<b>6</b>	<b>Solution .....</b>	<b>21</b>
6.1.	Normalization .....	21
6.2.	Model Training .....	25
6.3.	Result .....	27
6.3.1.	Accuracy Score .....	27
6.3.1.	Confusion Matrix .....	28
<b>7</b>	<b>Conclusion .....</b>	<b>30</b>

# 1 General Information

## Student Information

Student ID	Full Name	Email
20C14001	Le Duong Tuan Anh	<a href="mailto:leduongtuananh97@gmail.com">leduongtuananh97@gmail.com</a>

## Report Information

This report is a summary of the **Final Project** based on **Diabetes prediction with given Pima Indians medical details**.

## 2 Problem Statement

In this project, I use dataset is originally from the National Institute of Diabetes and Digestive and Kidney Diseases, which contains data between relationship between some medical details and the result on Diabetes.

**The objective** of the dataset is to diagnostically predict whether a patient has diabetes, based on certain diagnostic measurements included in the dataset. Several constraints were placed on the selection of these instances from a larger database. In particular, all patients here are females at least 21 years old of Pima Indian heritage.

The dataset consists of **several medical predictor variables** and **one target variable, Outcome**.

Predictor variables includes *the number of pregnancies the patient has had, their BMI, insulin level, age, and so on*.

### 3 Data Description

Total instances (records): 768

Total columns: 9 (8 medical attributes + 1 Outcome Class (has diabetes or not)).

Attribute's detail:

Column	Type	Datatype	Has missing value?
Pregnancies	Interval-scaled	Int64	No
Glucose	Interval-scaled	Int64	No
BloodPressure	Interval-scaled	Int64	No
SkinThickness	Interval-scaled	Int64	No
Insulin	Interval-scaled	Int64	No
BMI	Interval-scaled	Float	No
DiabetesPedigreeFunction	Interval-scaled	Float	No
Age	Interval-scaled	Int64	No
Outcome	Categorical Data Binary (0/1)	Int64	No

### Attribute's meaning:

**Pregnancies:** Number of times pregnant

**Glucose:** Plasma Glucose Concentration.

**BloodPressure:** Diastolic Blood Pressure.

**SkinThickness:** Estimate body fat.

**Insulin:** 2-Hour Serum Insulin.

**BMI:** Body Mass Index.

**DiabetesPedigreeFunction:** information about diabetes history in relatives and genetics.

**Age:** Age (years).

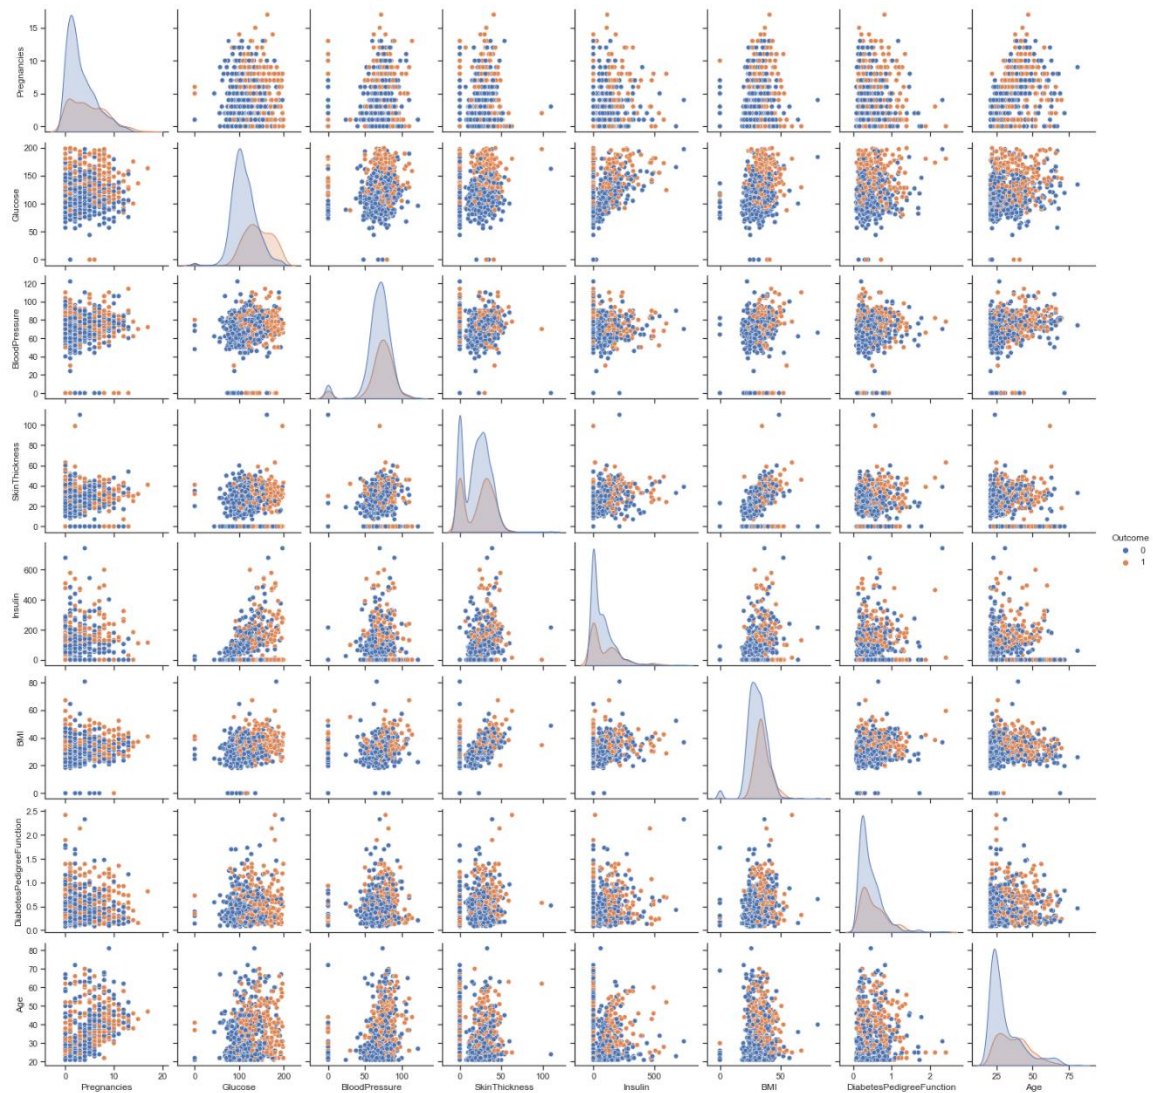
**Outcome:** 0 = Diabetic, 1 = Not Diabetic

# 4 Data Understanding

## 4.1 Description

	Pregnancies	Glucose	BloodPressure	SkinThickness	Insulin	BMI	DiabetesPedigreeFunction	Age	Outcome
<b>count</b>	768.000000	768.000000	768.000000	768.000000	768.000000	768.000000	768.000000	768.000000	768.000000
<b>mean</b>	3.845052	120.894531	69.105469	20.536458	79.799479	31.992578	0.471876	33.240885	0.348958
<b>std</b>	3.369578	31.972618	19.355807	15.952218	115.244002	7.884160	0.331329	11.760232	0.476951
<b>min</b>	0.000000	0.000000	0.000000	0.000000	0.000000	0.000000	0.078000	21.000000	0.000000
<b>25%</b>	1.000000	99.000000	62.000000	0.000000	0.000000	27.300000	0.243750	24.000000	0.000000
<b>50%</b>	3.000000	117.000000	72.000000	23.000000	30.500000	32.000000	0.372500	29.000000	0.000000
<b>75%</b>	6.000000	140.250000	80.000000	32.000000	127.250000	36.600000	0.626250	41.000000	1.000000
<b>max</b>	17.000000	199.000000	122.000000	99.000000	846.000000	67.100000	2.420000	81.000000	1.000000

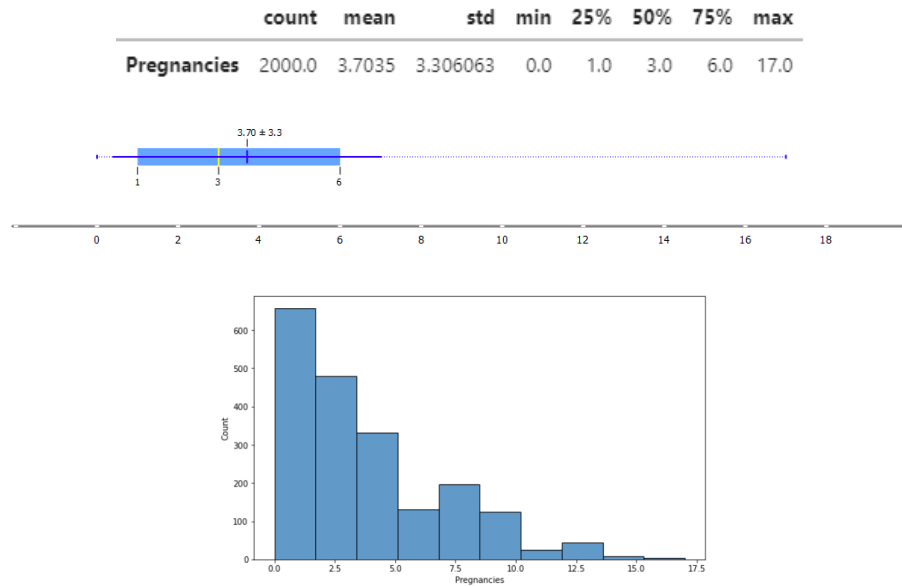
## 4.2 Overview by ScatterPlot






## 4.3 Statistical Information

### 2.1. Pregnancies





### 2.2. Glucose





## 2.3. BloodPressure

	count	mean	std	min	25%	50%	75%	max
BloodPressure	2000.0	69.1455	19.188315	0.0	63.5	72.0	80.0	122.0




## 2.4. SkinThickness

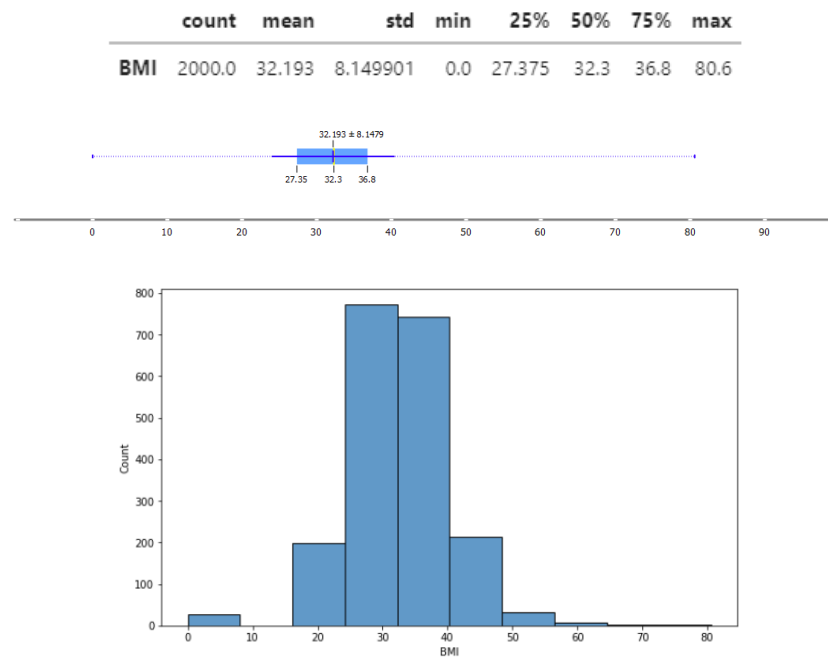
	count	mean	std	min	25%	50%	75%	max
SkinThickness	2000.0	20.935	16.103243	0.0	0.0	23.0	32.0	110.0




## 2.5. Insulin



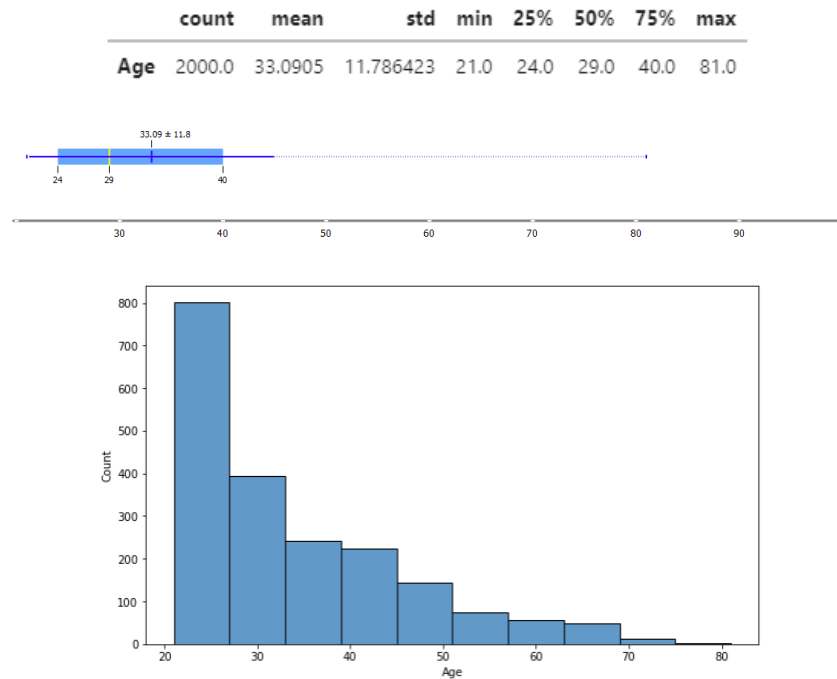
## 2.6. BMI



## 2.7. DiabetesPedigreeFunction




## 2.8. Age



## 2.9. Outcome

Outcome	
0	1316
1	684



## 5 Data Preprocessing

As we don't have any null-values, so we will check on zero-values and remove outliers (if needed).

### 5.1. Filling Zero-Values

column	total_zeros_value
Pregnancies	111
Glucose	5
BloodPressure	35
SkinThickness	227
Insulin	374
BMI	11
DiabetesPedigreeFunction	0
Age	0
Outcome	500

The **Pregnancies** column can be zero, as the gender might be male. There are 5 columns has zeros values: **Glucose, BloodPressure, SkinThickness, Insulin and BMI**. For example, SkinThickness cannot be zero, and BMI also. The other medical details cannot be zero, at least there must be very small values.

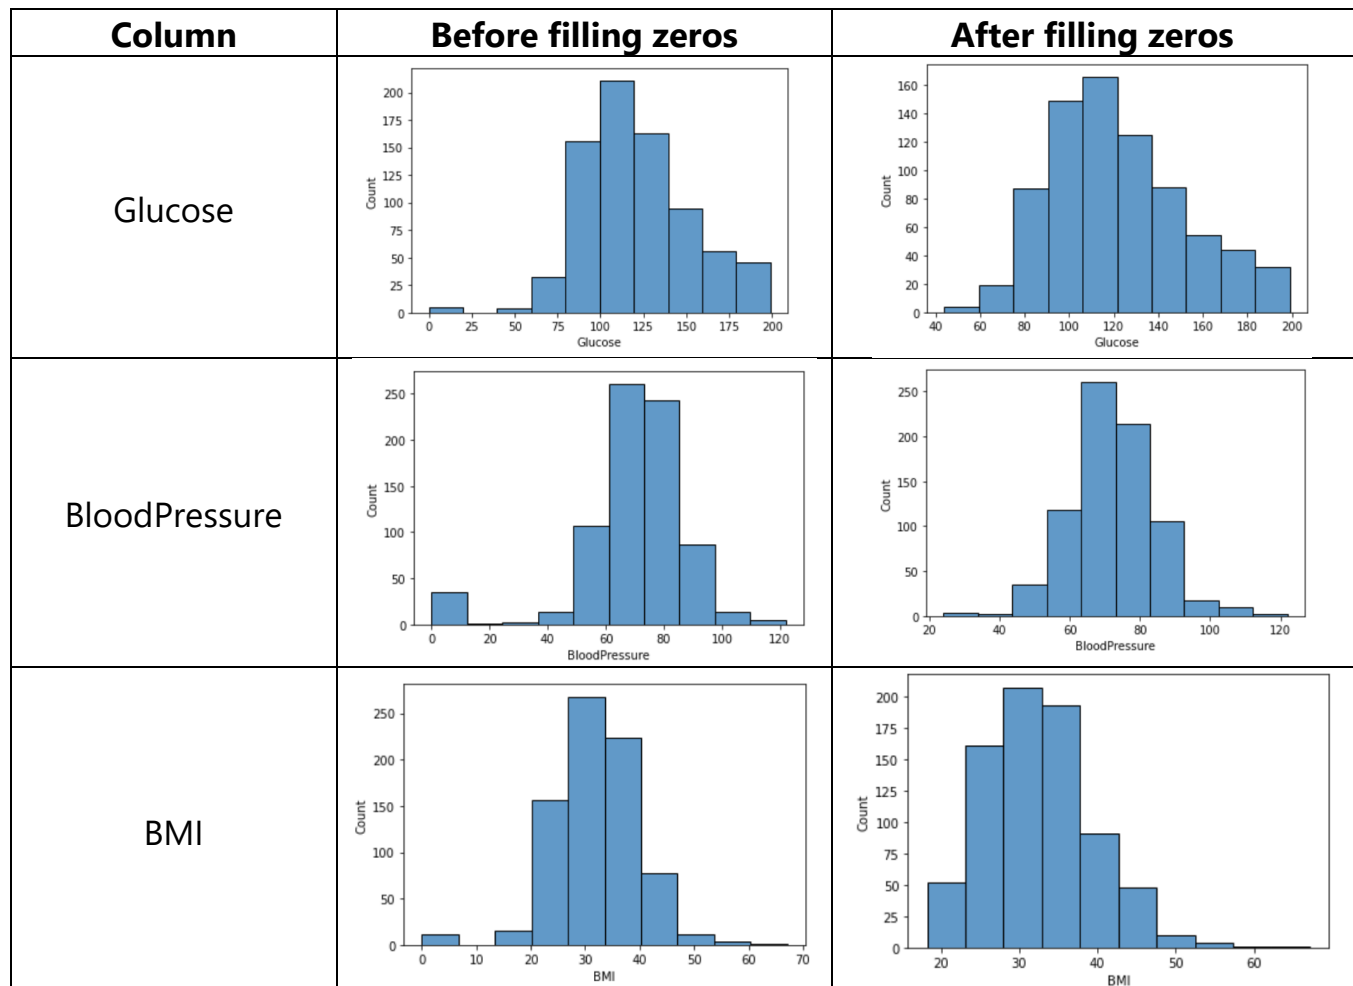
So, let's apply different type of zero-filling methods here.

### 5.1.1. Fill by Mean method

We just fill by mean on columns that have not much zero-values.

**The applied columns: Glucose, BloodPressure, BMI.**


```
# Glucose, BloodPressure, BMI can't be zero as well
# Let's fill by mean
for col in ['Glucose', 'BloodPressure', 'BMI']:
    data[col] = data[col].replace(value=data[col].mean(), to_replace=0)
# SkinThickness and Insulin should not have the zero values,
# but many are missing, so we skip at this step
```



### 5.1.2. Fill by Regression Method

The regression method is a very useful when the columns that having zero-values are very correlated with the other columns. So we will try this method on these columns: **Insulin**, **SkinThickness**.

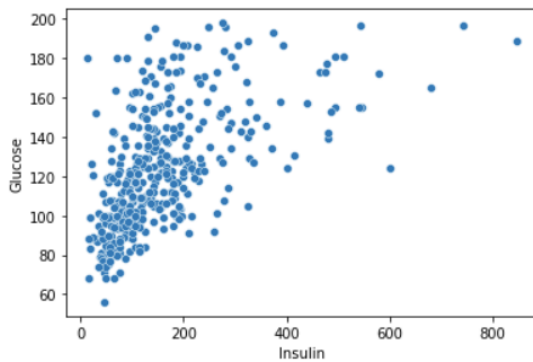
Correlation Matrix, **Pearson metric** (on dataset that have *Insulin*  $\neq 0$  and *SkinThickness*  $\neq 0$ ):






We can see that there are some high correlations between:

Insulin and Glucose (0.58).





SkinThickness and BMI (0.66).



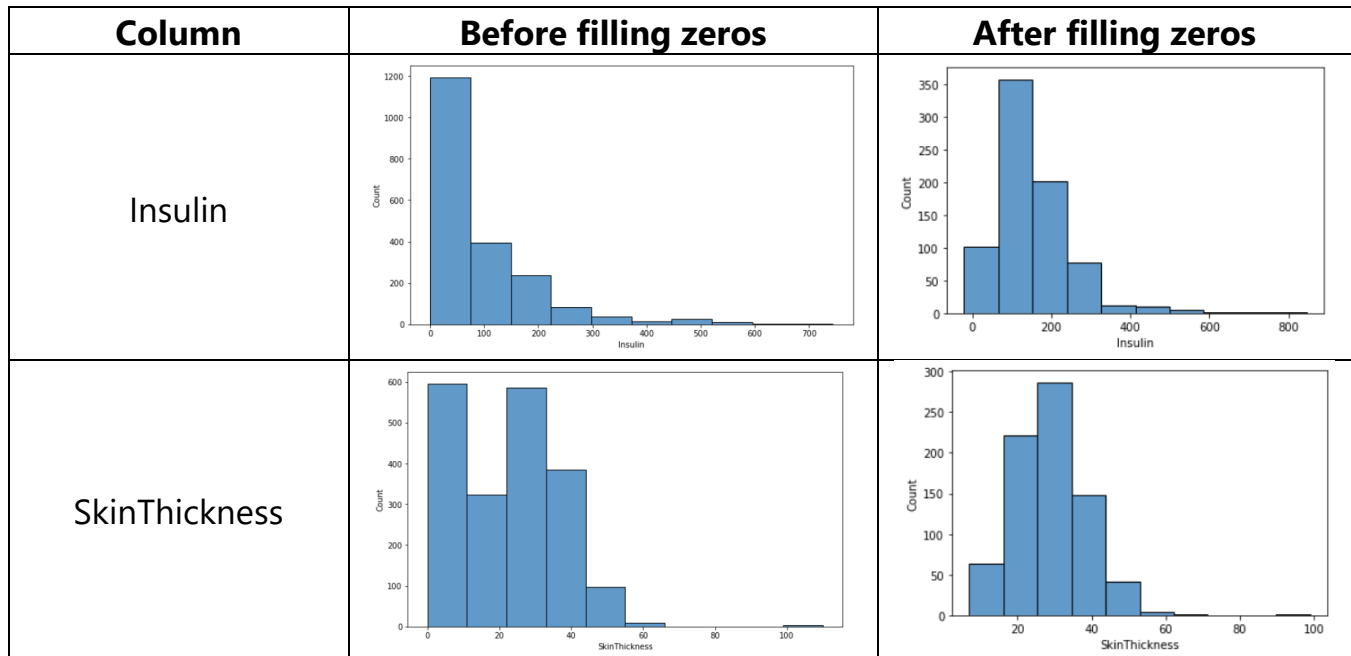
We try to estimate the zero-values by Linear Regression:

$$\text{Insulin} = \text{Glucose} * 2.23952761$$

$$\text{SkinThickness} = \text{BMI} * 0.99549695$$



Now the filling result:

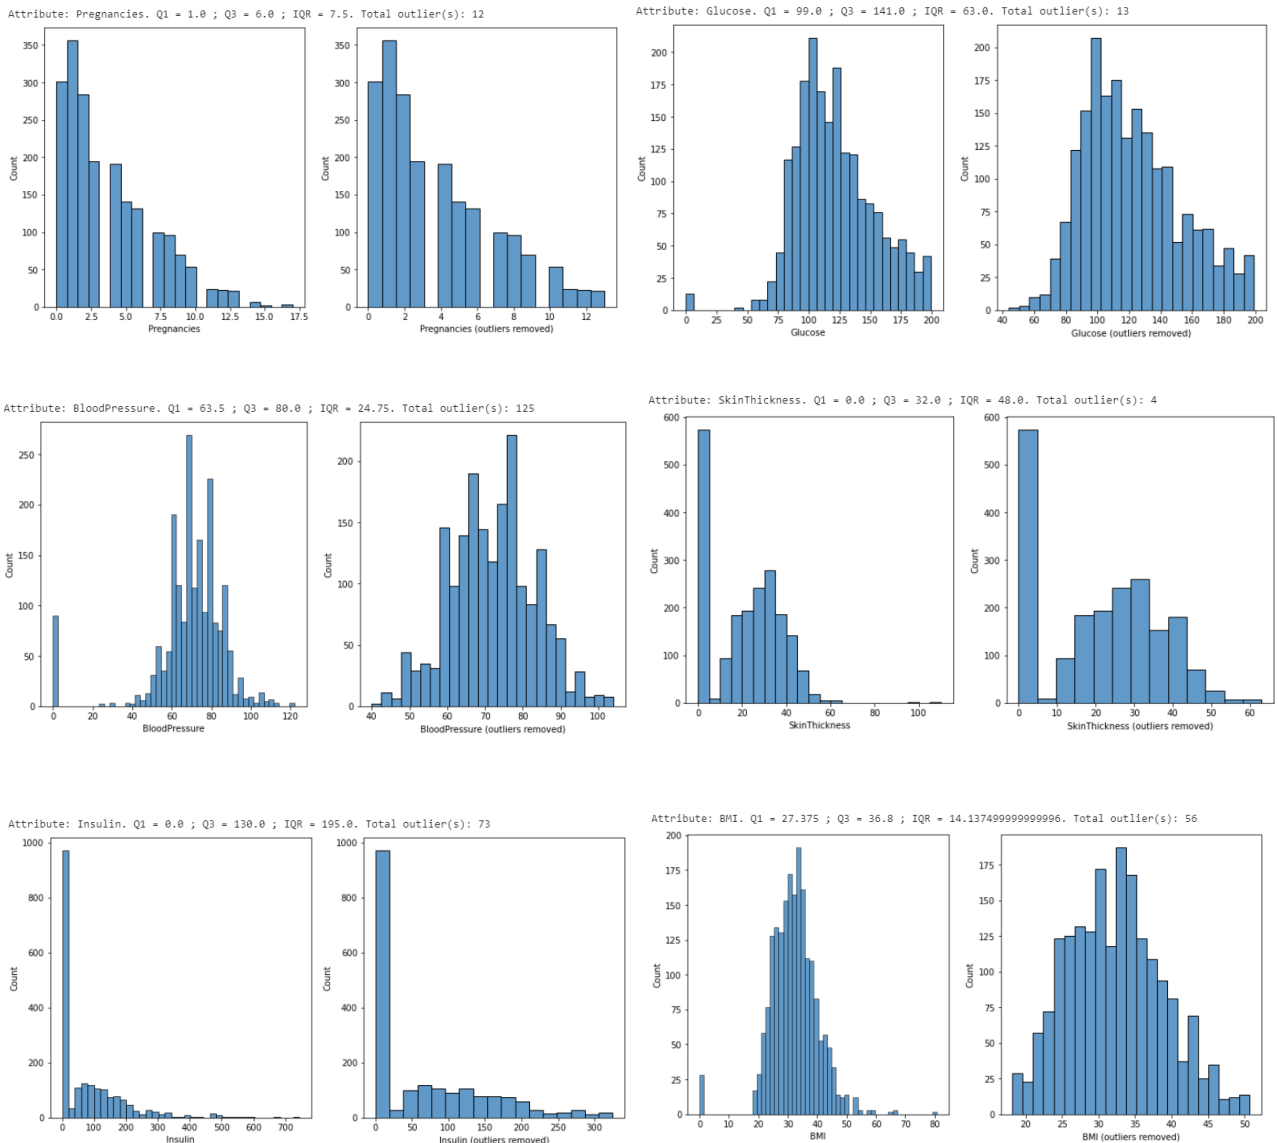


## 5.2. Outliers Removal



We would consider outliers by using **Inter-quartile Range** method. In case the data point is out of range

$$[ Q1 - 1.5 \cdot IQR, Q3 + 1.5 \cdot IQR ]$$



it might be **outliers**.



Attribute: DiabetesPedigreeFunction. Q1 = 0.244 ; Q3 = 0.624 ; IQR = 0.5700000000000001. Total outlier(s): 68



Attribute: Age. Q1 = 24.0 ; Q3 = 40.0 ; IQR = 24.0. Total outlier(s): 48



**Total removed outliers in dataset: 11,2%**

## 6 Solution

In this section, I will do some different classification methods to classify if the people have diabetes or not, based on their medical details.

Before we do the classification tasks, we need to normalize the input first.

### 6.1. Normalization

We standardize features by removing the mean, then scaling dataset to unit variance.

Note that we will not standardize the **Outcome** column, as it is a **categorical column**.

The standardized score is calculated as below:

$$Z = (X - U) / S$$

**Where:**

- X: The value of sample in dataset.
- U: The mean of dataset.
- S: The standard deviation of dataset.

```
df = data.drop(['Outcome'], axis=1)
```


```
df.mean()
```


```
Pregnancies      3.70350  
Glucose          121.18250  
BloodPressure    69.14550  
SkinThickness    20.93500  
Insulin          80.25400  
BMI             32.19300  
DiabetesPedigreeFunction  0.47093  
Age             33.09050  
dtype: float64
```

```
df.std()
```

```
Pregnancies      3.306063  
Glucose          32.068636  
BloodPressure    19.188315  
SkinThickness    16.103243  
Insulin          111.180534  
BMI             8.149901  
DiabetesPedigreeFunction  0.323553  
Age             11.786423  
dtype: float64
```

The result after normalization (*the blue shows original, the red shows normalized*)







## 6.2. Model Training

In this project, I have done on some different kind of classification methods as below:

- Logistic Regression

[https://scikit-learn.org/stable/modules/generated/sklearn.linear\\_model.LogisticRegression.html](https://scikit-learn.org/stable/modules/generated/sklearn.linear_model.LogisticRegression.html)

- Decision Tree

<https://scikit-learn.org/stable/modules/generated/sklearn.tree.DecisionTreeClassifier.html#sklearn.tree.DecisionTreeClassifier>

- Random Forest

<https://scikit-learn.org/stable/modules/generated/sklearn.ensemble.RandomForestClassifier.html>

- SVM (Linear Kernel)

<https://scikit-learn.org/stable/modules/generated/sklearn.svm.SVC.html>

- kNN

<https://scikit-learn.org/stable/modules/generated/sklearn.neighbors.KNeighborsClassifier.html>

The implementation has been done by scikit-learn library (Python), so I just use the model in the library.

The data is split to **trainset** and **testset**, with ratio **7/3**: 70% data for training and 30% for testing, respectively.

*The splitting data and training code:*

```
: # We need to split dataset to train/test data
# 70% train, 30% test
from sklearn.model_selection import train_test_split
x_train, x_test, y_train, y_test = train_test_split(x, y, train_size=0.7)

: # We will do some different type of algorithms to determine which is the best
from sklearn.linear_model import LogisticRegression
from sklearn.tree import DecisionTreeClassifier
from sklearn.ensemble import RandomForestClassifier
from sklearn.svm import SVC
from sklearn.neighbors import KNeighborsClassifier

list_models = [
    LogisticRegression(),
    DecisionTreeClassifier(),
    RandomForestClassifier(),
    SVC(),
    KNeighborsClassifier(),
]


: # Now training
for model in list_models:
    model.fit(x_train.values, y_train)
```

## 6.3. Result

The testing result are calculated on 30% data (testset).

### 6.3.1. Accuracy Score

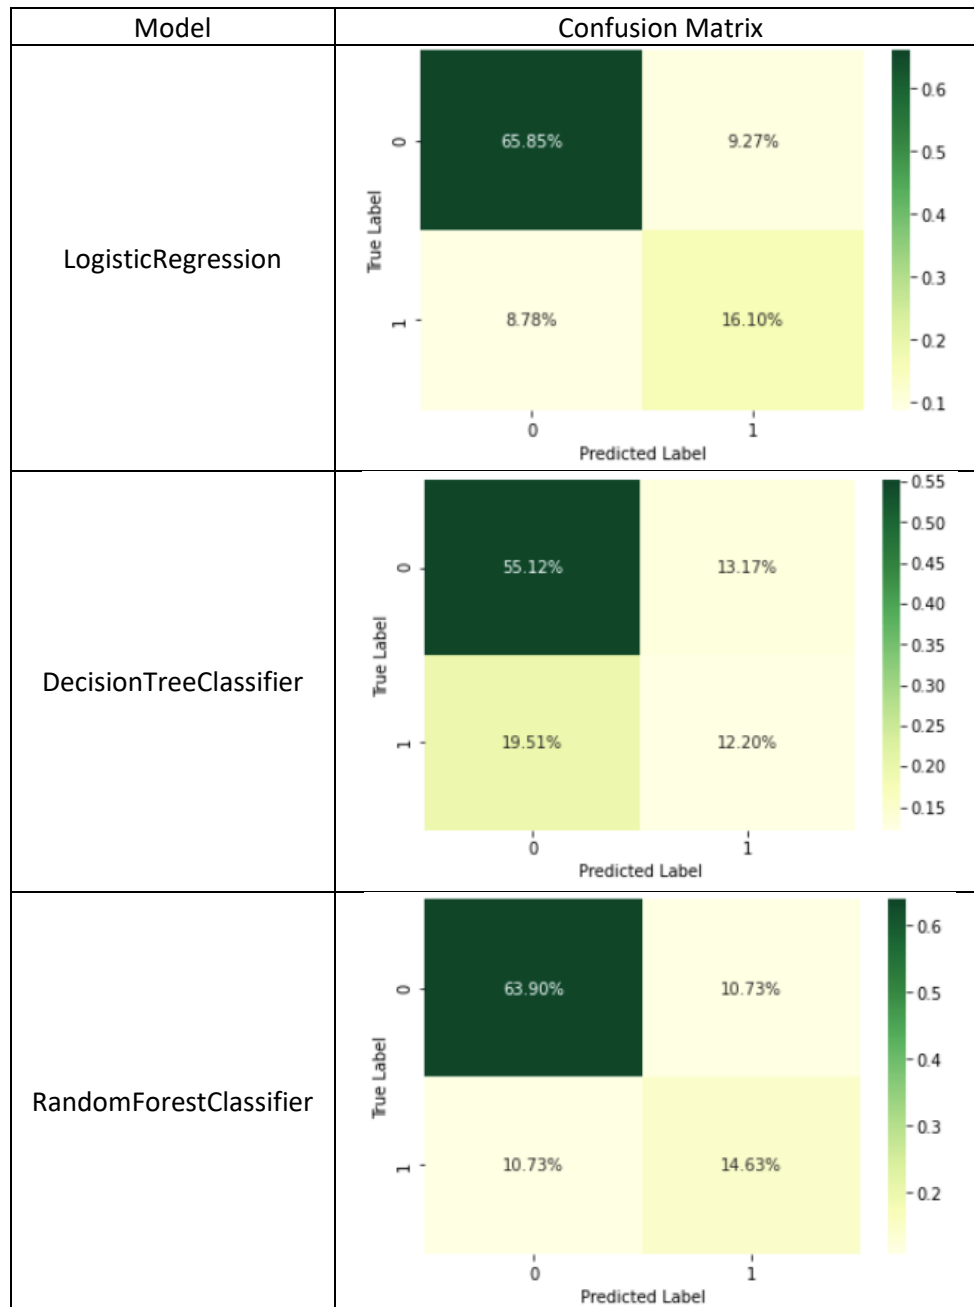
model_name	accuracy_score
LogisticRegression	0.819512
DecisionTreeClassifier	0.673171
RandomForestClassifier	0.785366
SVC	0.790244
KNeighborsClassifier	0.756098

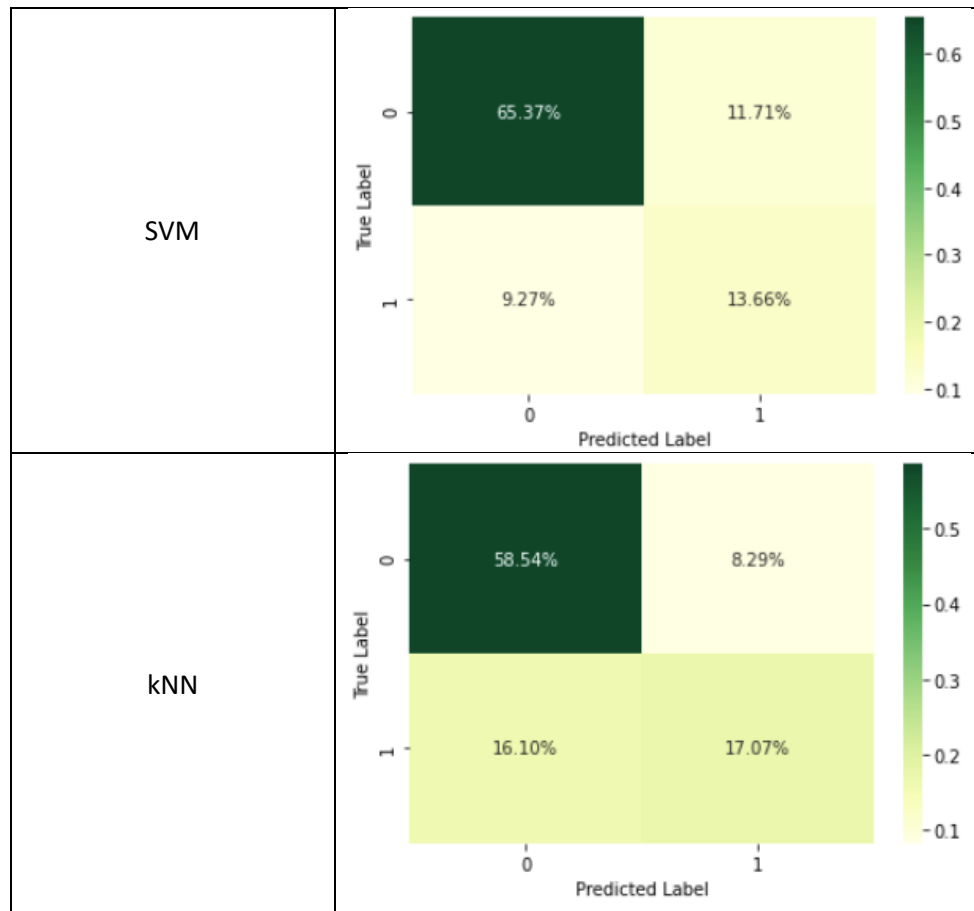


With accuracy metrics, the **LogisticRegression** gave the best result.

### 6.3.1. Confusion Matrix

To analyze more in detail on the prediction of classification problems, we use confusion matrix.





This classification is done on **medical examination**, which is the False Negative is one of the most importance metrics. Based on this result, Logistic Regression method gave the best result with **8.78% (lower is better)**.

## 7 Conclusion

In this project, I have done end-to-end flow to analyze, process and do some classification methods on the Diabetes Prediction task.

The precision, recall and f-score of all methods are below:

model_name	precision	recall	fscore
LogisticRegression	0.758484	0.761841	0.760128
DecisionTreeClassifier	0.609666	0.595879	0.599341
RandomForestClassifier	0.716566	0.716566	0.716566
SVC	0.707139	0.721923	0.713696
KNeighborsClassifier	0.728695	0.695309	0.705460

In conclusion, **the simple Logistic Regression** performs the **best result** on the provided dataset.

This result might be improved if we use other complex methods, e.g., Deep Learning.